

```
In [ ]: # Simple Regression Using "MachineHoursCurrentMeter" to Calculate "SalePrice"
```

```
In [1]: import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
```


```
In [ ]: data = pd.read_csv('fast iron 100k data.csv')
```

```
In [3]: data
```

```
Out[3]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCur
0	1139283	36000	1052214	2232	121	3	1998	
1	1139299	9500	1002713	21442	121	3	2003	
2	1139356	53000	1004127	25458	121	3	2000	
3	1139369	14000	1050658	1918	121	3	1000	
4	1139390	8250	1011114	7110	121	3	1000	
...
99995	1417929	93500	26366	3357	132	6	1990	
99996	1417930	47000	1412330	3357	132	9	1990	
99997	1417931	47000	1412330	3357	132	9	1990	
99998	1417934	48000	1412330	3357	132	9	1990	
99999	1417937	92500	1195271	3357	132	1	1990	

100000 rows × 53 columns



```
In [4]: data['MachineHoursCurrentMeter'] = data['MachineHoursCurrentMeter'].replace(np.nan, 0)
```

```
In [5]: data.describe()
```

```
Out[5]:
```


	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Year
count	1.000000e+05	100000.00000	1.000000e+05	100000.000000	100000.000000	100000.000000	100000.0
mean	1.294013e+06	32361.55013	1.067671e+06	5485.211650	129.362860	6.502580	1931.4
std	7.910623e+04	23797.71391	3.947570e+05	5078.726263	4.696196	15.605157	237.2
min	1.139246e+06	4750.00000	1.300000e+01	28.000000	121.000000	1.000000	1000.0
25%	1.252166e+06	15500.00000	1.004246e+06	3171.000000	132.000000	1.000000	1984.0
50%	1.292648e+06	24500.00000	1.151534e+06	3819.000000	132.000000	3.000000	1993.0
75%	1.363022e+06	41000.00000	1.355657e+06	6797.000000	132.000000	4.000000	1999.0
max	1.417937e+06	142000.00000	1.558285e+06	37198.000000	132.000000	99.000000	2010.0

```
In [6]:
```

```
y = data['SalePrice']
x1 = data['MachineHoursCurrentMeter']
```

```
In [7]:
```

```
plt.scatter(x1, y)
plt.xlabel('SalePrice', fontsize = 20)
plt.ylabel('MachineHoursCurrentMeter', fontsize = 20)
plt.show()
```



```
In [8]:
```

```
x = sm.add_constant(x1)
results = sm.OLS(y, x).fit()
results.summary()
```

```
Out[8]:
```

OLS Regression Results

Dep. Variable:	SalePrice	R-squared:	0.012
Model:	OLS	Adj. R-squared:	0.012
Method:	Least Squares	F-statistic:	1221.

Date: Thu, 21 Dec 2023 **Prob (F-statistic):** 7.90e-266
Time: 10:19:24 **Log-Likelihood:** -1.1490e+06
No. Observations: 100000 **AIC:** 2.298e+06
Df Residuals: 99998 **BIC:** 2.298e+06
Df Model: 1
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	3.15e+04	78.754	399.985	0.000	3.13e+04	3.17e+04
MachineHoursCurrentMeter	0.7231	0.021	34.938	0.000	0.683	0.764

Omnibus: 26748.274 **Durbin-Watson:** 0.815
Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 60472.371
Skew: 1.539 **Prob(JB):** 0.00
Kurtosis: 5.245 **Cond. No.** 4.01e+03

Notes:


- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.01e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [9]:

```

plt.scatter(x1, y)
yhat = 0.7231 * x1 + 31500 # Regression Equation
fig = plt.plot(x1, yhat, lw = 4, c = 'orange', label = 'regression line')
plt.xlabel('SalePrice', fontsize = 20)
plt.ylabel('MachineHoursCurrentMeter', fontsize = 20)
plt.show()

```



```
In [10]: print(results.summary()) # Notice the really Low R-squared
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          SalePrice      R-squared:                0.012
Model:                  OLS           Adj. R-squared:            0.012
Method:                 Least Squares  F-statistic:              1221.
Date:                  Thu, 21 Dec 2023  Prob (F-statistic):       7.90e-266
Time:                  10:19:29        Log-Likelihood:           -1.1490e+06
No. Observations:      100000         AIC:                    2.298e+06
Df Residuals:          99998         BIC:                    2.298e+06
Df Model:               1
Covariance Type:       nonrobust
=====
=====
                        coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
const                  3.15e+04    78.754    399.985    0.000    3.13e+04    3.17
e+04
MachineHoursCurrentMeter  0.7231    0.021    34.938    0.000    0.683
0.764
=====
Omnibus:                26748.274    Durbin-Watson:           0.815
Prob(Omnibus):           0.000    Jarque-Bera (JB):        60472.371
Skew:                    1.539    Prob(JB):                0.00
Kurtosis:                5.245    Cond. No.                4.01e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.01e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [ ]: # Multiple Regression Using "MachineHoursCurrentMeter" and "YearMade" to Calculate "SalePrice"
```

```
In [11]: import numpy as np
import pandas as pd
```

```
import matplotlib.pyplot as plt
import statsmodels.api as sm
```


```
In [ ]: data = pd.read_csv('fast iron 100k data.csv')
```

```
In [13]: data.head()
```

```
Out[13]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentM
0	1139283	36000	1052214	2232	121	3	1998	
1	1139299	9500	1002713	21442	121	3	2003	
2	1139356	53000	1004127	25458	121	3	2000	
3	1139369	14000	1050658	1918	121	3	1000	
4	1139390	8250	1011114	7110	121	3	1000	

5 rows × 53 columns



```
In [14]: data.describe()
```

```
Out[14]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Year
count	1.000000e+05	100000.00000	1.000000e+05	100000.000000	100000.000000	100000.000000	100000.0
mean	1.294013e+06	32361.55013	1.067671e+06	5485.211650	129.362860	6.502580	1931.4
std	7.910623e+04	23797.71391	3.947570e+05	5078.726263	4.696196	15.605157	237.2
min	1.139246e+06	4750.00000	1.300000e+01	28.000000	121.000000	1.000000	1000.0
25%	1.252166e+06	15500.00000	1.004246e+06	3171.000000	132.000000	1.000000	1984.0
50%	1.292648e+06	24500.00000	1.151534e+06	3819.000000	132.000000	3.000000	1993.0
75%	1.363022e+06	41000.00000	1.355657e+06	6797.000000	132.000000	4.000000	1999.0
max	1.417937e+06	142000.00000	1.558285e+06	37198.000000	132.000000	99.000000	2010.0



```
In [15]: data['MachineHoursCurrentMeter'] = data['MachineHoursCurrentMeter'].replace(np.nan, 0)
data = data[data['YearMade'] != 1000] # Excludes rows in which YearMade has a value of
```

```
In [16]: data.describe() # Notice how there are now only 93,916 rows
```

Out[16]:		SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade
	count	9.391600e+04	93916.000000	9.391600e+04	93916.000000	93916.000000	93916.000000	93916.0000
	mean	1.299405e+06	32842.654372	1.068873e+06	5517.963989	129.695664	6.599664	1991.7818
	std	7.780581e+04	24016.172084	4.013234e+05	5121.838352	4.476376	15.859059	9.7477
	min	1.139246e+06	4750.000000	1.300000e+01	28.000000	121.000000	1.000000	1919.0000
	25%	1.258933e+06	16000.000000	1.001703e+06	3171.000000	132.000000	1.000000	1986.0000
	50%	1.298802e+06	25000.000000	1.164262e+06	3819.000000	132.000000	3.000000	1994.0000
	75%	1.366180e+06	42000.000000	1.362241e+06	6805.000000	132.000000	4.000000	1999.0000
	max	1.417937e+06	142000.000000	1.558285e+06	37198.000000	132.000000	99.000000	2010.0000



```
In [17]: y = data['SalePrice']
x1 = data[['MachineHoursCurrentMeter', 'YearMade']]
```

```
In [18]: x = sm.add_constant(x1)
results = sm.OLS(y, x).fit()
results.summary()
```

Out[18]:

OLS Regression Results							
Dep. Variable:	SalePrice	R-squared:	0.054				
Model:	OLS	Adj. R-squared:	0.054				
Method:	Least Squares	F-statistic:	2673.				
Date:	Thu, 21 Dec 2023	Prob (F-statistic):	0.00				
Time:	10:19:54	Log-Likelihood:	-1.0779e+06				
No. Observations:	93916	AIC:	2.156e+06				
Df Residuals:	93913	BIC:	2.156e+06				
Df Model:	2						
Covariance Type:	nonrobust						
		coef	std err	t	P> t 	[0.025	0.975]
	const	-9.593e+05	1.61e+04	-59.761	0.000	-9.91e+05	-9.28e+05
MachineHoursCurrentMeter		0.6337	0.028	22.951	0.000	0.580	0.688
	YearMade	497.8320	8.062	61.747	0.000	482.030	513.634
Omnibus:	22790.243	Durbin-Watson:	0.786				
Prob(Omnibus):	0.000	Jarque-Bera (JB):	47872.793				
Skew:	1.435	Prob(JB):	0.00				
Kurtosis:	5.000	Cond. No.	6.52e+05				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.52e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [19]:

```
print(results.summary()) # R-squared increased from to what it was in simple regression
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          SalePrice      R-squared:                0.054
Model:                  OLS           Adj. R-squared:            0.054
Method:                 Least Squares   F-statistic:              2673.
Date:                   Thu, 21 Dec 2023   Prob (F-statistic):       0.00
Time:                   10:19:56          Log-Likelihood:           -1.0779e+06
No. Observations:       93916            AIC:                     2.156e+06
Df Residuals:           93913            BIC:                     2.156e+06
Df Model:                2
Covariance Type:        nonrobust
=====
=====
=====
coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
const      -9.593e+05    1.61e+04   -59.761    0.000    -9.91e+05    -9.28
e+05
MachineHoursCurrentMeter    0.6337      0.028     22.951    0.000      0.580
0.688
YearMade    497.8320      8.062     61.747    0.000     482.030     51
3.634
=====
Omnibus:            22790.243    Durbin-Watson:           0.786
Prob(Omnibus):      0.000    Jarque-Bera (JB):        47872.793
Skew:               1.435    Prob(JB):                0.00
Kurtosis:           5.000    Cond. No.                 6.52e+05
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.52e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In []:

```
# Regression Equation: SalePrice = 0.6337 * MachineHoursCurrentMeter + 497.8320 * YearMade
```